

Systems Design and Security



Part 9: Java and MySQL

<http://staffwww.dcs.shef.ac.uk/people/A.Simons/>

Home ⇒ Teaching ⇒ Lectures
⇒ COM2008/COM3008



Bibliography

- Database Systems

- T Connolly and C Begg, Database Systems – a Practical Approach to Design, Implementation and Management, 6th ed., Pearson, 2014.
- C J Date, An Introduction to Database Systems, 8th ed., Pearson, 2003.

- MySQL and Java

- K Sharan, Beginning Java 8 APIs, Extensions and Libraries, APress, 2014. -- Java 8.
- M Matthews, J Cole, J D Gradecki, MySQL and Java Developers Guide, Wiley, 2003. -- Java 1.4
- take care with older Java versions




Outline

- Installing MySQL and JDBC
- Managing connection resources
- Executing updates and queries
- Object-data conversion
- Commits and transactions
- Injection attack and validation

Programs and Databases

- Weakly-typed databases
 - databases store mostly text, some binary data
 - SQL uses simple types: int, char, date, etc.
 - search results mostly strings, chars, ints
- Strongly-typed programming languages
 - programs use rich, structured object types
 - stronger type-checking at the class-level
- Programmatic access to data
 - programs have to overcome the “impedance mismatch”
 - map string, int data into complex object types
 - vice-versa, when storing objects back in databases





MySQL Database

- MySQL is a robust free database
 - runs under Windows, Unix and Linux
 - uses the standard SQL query language
 - many on-line tutorials are available
- Get a MySQL group account (see later)
 - your lecturer will tell you when these are ready
 - all your group members get shared access to the same DB
 - only accessible from campus network (or use VPN)



DCS Guide: <https://guide.dcs.shef.ac.uk/doku.php?id=students>
MySQL Guide, e.g.: <http://www.mysqltutorial.org/>



Java Database Connectivity


- JDBC developed by Sun Microsystems
 - Java 8 SE comes with JDBC 4.2
 - builds on the simplicity/portability of SQL
 - supports conversion to/from Java objects
 - database failures reported as Java exceptions
- JDBC is the natural API for Java access to data
 - simple to use for non-database programmers
 - hides specific database details from programmers
 - can access any tabular data source, eg: CSV files






JDBC Architecture

- Assumes multiple vendor RDB implementations
- Assumes SQL used as the common query language




JDBC Architecture

- JDBC defines a set of interfaces
- vendors implement these for their database engine
- known as a **JDBC driver**
- need to download a **specific** JDBC driver for each database engine



JDBC Architecture

- run Java apps, that use the JDBC interface
- easy to switch between DB vendors, hardly altering a line of Java code
- just load a different JDBC driver at Java startup
- if you start with toy DBs in Microsoft Access...
- ... can scale up later to large Oracle DBs





Overview

Platform-independent
Connector/J in your
group project folder

- Get the JDBC driver for MySQL
 - from: <http://www.mysql.com/products/connector/>
 - use the **Connector/J** version of the driver (for Java)
 - install the **jarfile** on your project's build path
 - see: <http://download.oracle.com/javase/tutorial/jdbc/>
- Find/read about the JDBC classes
 - the packages: **java.sql**, **javax.sql** contain JDBC classes
 - **DriverManager** – creates connections to the DB
 - **Connection** – represents an open connection to the DB
 - **Statement** – represents a statement to execute on the DB
 - **ResultSet** – is an iterator over the results of a query



Install the Driver

- Unzip the download bundle
 - the download is: `mysql-connector-java-[version].zip`
 - contains a jarfile: `mysql-connector-java-[version]-bin.jar`
 - this contains a driver class: `com.mysql.jdbc.Driver`
- Java archive (jar) files
 - zipped files containing compiled Java classes, docs, etc.
 - can contain library of add-on software (as in this case)
 - need to tell Java **where to find** the jarfiles you use
 - JDK: append path to the jarfile to your **CLASSPATH**
 - Eclipse, IntelliJ, NetBeans: add the jarfile to your **build path**



Eclipse Build Path

- Select project in the package explorer list
 - Right-click to pop up main menu
 - Select **Build Path** and submenu **Configure Build Path**
- Properties Browser opens at **Java Build Path**
 - Select tab **Libraries**
 - Click on right-hand button **Add External JAR**
 - Browse for **mysql-connector-java-[version]-bin.jar**
 - Open (**OK**), it should then appear in the list
 - Click on **OK** to save and close Properties Browser.



IntelliJ and Netbeans

■ IntelliJ IDEA

- Select project in the project list
- Go to main **File** menu → **Project Structure**
- Go to **Project Settings** → **Libraries**
- Click **+** icon to add external library, browse for JAR, OK

■ NetBeans

- Right-click on the project for main menu
- Select **Properties** → **Libraries**
- Right-click **Libraries** → **Add JAR/Folder**
- Under **Compile** tab, browse for JAR, OK



Check Driver Exists

```
import java.sql.*;
import java.util.*;

public class FindDrivers {
    public static void main(String[] args) throws Exception {
        System.out.println("\nDrivers loaded as properties:");
        System.out.println(System.getProperty("jdbc.drivers"));
        System.out.println("\nDrivers loaded by DriverManager:");
        Enumeration<Driver> list = DriverManager.getDrivers();
        while (list.hasMoreElements())
            System.out.println(list.nextElement());
    }
}
```



Example Output

```
import java.sql.*;
import java.util.*;

public class FindDrivers {
    public static void main(String[] args) throws Exception {
        System.out.println("\nDrivers loaded as properties:");
        System.out.println(System.getProperty("jdbc.drivers"));
        System.out.println("\nDrivers loaded by DriverManager:");
        EnumDrivers loadedAsProperties = DriverManager.getDrivers();
        while (loadedAsProperties.hasMoreElements()) {
            System.out.println(loadedAsProperties.nextElement());
        }
    }
}
```

Drivers loaded as properties:

Drivers loaded by DriverManager:

sun.jdbc.odbc.JdbcOdbcDriver@ca0b6

com.mysql.jdbc.Driver@1270b73



Using JDBC

- Use the JDBC 4.x library
 - `import java.sql.*;` // adequate for most purposes
 - `import javax.sql.*;` // only for advanced features
- Four main steps
 - open a connection to the server
 - execute a query/update (one or more)
 - iterate over the results of a query
 - release server resources - **otherwise will crash!**
 - typically after 20 opened connections!
- Resource management
 - need to get this right, remember to close (see next)

Resources – Old Style

```
Connection con = null; // connection to a database
```

```
try {  
    con = DriverManager.getConnection(...);
```

declare *con* outside
the try-block

```
    // use the open connection  
    // for one or more queries
```

opening the
connection may fail

```
}  
catch (Exception ex) {  
    ex.printStackTrace();  
}
```

handle any exception
and report the error

```
finally {  
    if (con != null) con.close();  
}
```

finally, close *con* if it
was opened

- Explicit release of resources, in a `try...catch...finally` block

Resources – New Style

```
try (Connection con = DriverManager.getConnection(...)) {
```

```
    // use the open connection  
    // for one or more queries
```

```
}  
catch (Exception ex) {  
    ex.printStackTrace();  
}
```

declare *con* inside ()
parentheses, before the
start of the try-block

automatically closed at
the end of the block

- Implicit release of resources, in a **try-with-resources** block
- Any **AutoCloseable** object can be initialised in the () parentheses
- Automatically closed, whether success or failure occurs
- Since Java 7 / JDK1.7, **if your IDE recognises it!**

<https://docs.oracle.com/javase/tutorial/essential/exceptions/tryResourceClose.html>

Open a Connection

```
Connection con = null; // a Connection object

try {

    con = DriverManager.getConnection(
        "jdbc:mysql://server/dbname", "userID", "password");

    // use the open connection
    // for several queries

}
catch (Exception ex) {
    ex.printStackTrace();
}
finally {
    if (con != null) con.close();
}
```

`server/dbname` path may include a port number, and `dbname` is as you set it up
`userID` and `password` are the user ID and password for the `dbname`

All other query code goes in here

Different Styles

- Example database IDs
 - on separate DB server, with DB name "team043"
 - on localhost, on a specific port, with DB name "myDB"

```
"jdbc:mysql://stusql.dcs.shef.ac.uk/team043"
```

```
"jdbc:mysql://localhost:3306/myDB"
```

- Example connection methods
 - multi-argument, or single string URL argument

```
Connection con = DriverManager.getConnection(  
    "database", "dbuser", "dbpassword");
```

```
Connection con = DriverManager.getConnection(  
    "database?user=dbuser&password=dbpassword");
```

Execute an Update

```
Statement stmt = null;

try {
    stmt = con.createStatement();
    int count = stmt.executeUpdate(
        "UPDATE lecturer SET office = 119"
        + " WHERE name = 'A Simons'");
}
catch (SQLException ex) {
    ex.printStackTrace();
}
finally {
    if (stmt != null)
        stmt.close();
}
```

Use `executeQuery(sqlString)` for SELECT statements

Use `executeUpdate(sqlString)` for INSERT, DELETE and UPDATE statements

Updates return the number of rows that were updated (or zero)

Data Manipulation

Command	Action	Example
CREATE	creates a table	CREATE TABLE student (...);
INSERT	inserts records	INSERT INTO student VALUES (...);
UPDATE	modifies records	UPDATE student SET name='Jill Smith' WHERE id=4;
DELETE	deletes records	DELETE FROM student WHERE name='Joe Bloggs';
DESCRIBE	info about a table	DESCRIBE student;
SHOW	info about system	SHOW DATABASES; SHOW TABLES [FROM db];
DROP	delete tables	DROP TABLE student;

Execute a Query

```
Statement stmt = null; // a SQL statement object

try {
    stmt = con.createStatement();
    ResultSet res =
        stmt.executeQuery("SELECT * FROM lecturer");

    // do what you like with the results;
    // could convert into objects

    res.close();
}
catch (SQLException ex) {
    ex.printStackTrace();
}
finally {
    if (stmt != null) stmt.close();
}
```

Statement objects are created by the open Connection

executeQuery() accepts a SQL query and returns a ResultSet

Remember to close Statement objects when finished

A decorative graphic consisting of overlapping yellow, red, and blue squares with a black crosshair.

SQL Data Access

Command	Action	Example
SELECT	displays a table	SELECT * FROM student;
SELECT	projects columns	SELECT id, name FROM student;
WHERE	selects rows	SELECT id, name FROM student WHERE id=4;
WHERE	inner join	SELECT name FROM lecturer, module WHERE module.lec = lecturer.id
INNER JOIN	inner join	SELECT name FROM lecturer INNER JOIN module ON lec = lecturer.id
USING	natural join	SELECT name FROM module, student USING id

Iterate over Results

```
ResultSet res = stmt.executeQuery(  
    "SELECT * FROM lecturer WHERE name='A Simons'");  
while (res.next()) {  
    int id = res.getInt(1);           // col 1 as int  
    String name = res.getString(2);  // col 2 as string  
    int office = res.getInt(3);       // col 3 as int  
  
    Lecturer teacher = new Lecturer(id, name, office);  
  
    // do something with teacher  
    // eg: store object in a list  
}  
  
res.close();
```

ResultSet objects work like iterators
extract columns by integer index, or by
column name, supplied as a string
getX() for each of several Java types X

Result Set

- Provides access to results of DB queries one row at a time
- Is a reference to the actual data, one row at a time
- The actual set of results may be very large

```
ResultSet res = stmt.executeQuery(  
    "SELECT * FROM lecturer WHERE name='A Simons'");  
while (res.next()) {  
    int id = res.getInt("id");    // access by col name  
    String name = res.getString("name");  
    int office = res.getInt("office");  
  
    Lecturer teacher =  
        new Lecturer(id, name, office);  
  
    // do something with teacher  
}
```

id	name	office
15	A Simons	119
23	A Stratton	118



Data Conversion

```
class ResultSet {  
  
    public int getInt(int column);  
    public String getString(int columnIndex);  
    public float getFloat(int columnIndex);  
    public double getDouble(int columnIndex);  
    public Date getDate(int columnIndex);  
    public Time getTime(int columnIndex);  
    ...  
    public int getInt(String columnLabel);  
    public String getString(String columnLabel);  
    ...  
    public Object getObject(int columnIndex);  
    public Object getObject(String columnLabel);  
}
```

Use specific access method where the type is known

Access using a column index or column name

Use the getObject() method if type is unknown

Performs a default type conversion, result is in a var of type Object



Lab 1: Execute Query

Run a Poll



- Sketch all the Java for a query
 - connect to the DB
 - create the Statement
 - execute a query seeking all Lecturers whose first initial is 'A' (assume a single name-field)
 - iterate through the ResultSet
 - return the result as a List<Lecturer>
- Remember to manage resources
 - always close the ResultSet
 - always close the Statement
 - always close the Connection



When to Commit?

- Commit changes immediately
 - good strategy in most cases
 - saves important updates, as these happen
 - fewer concurrent data access issues
 - avoid dirty reads – reading uncommitted changes
 - avoid non-repeatable reads – rows altered while reading
 - avoid phantom reads – rows inserted while reading
- Defer committing changes
 - execute a batch of insert, update jobs for efficiency
 - group sets of changes that logically belong together
 - notion of transactions (see later)

Control Auto-Commit

- Default JDBC setting
 - by default, every `executeUpdate()` commits immediately to the database
- Custom JDBC setting
 - can control the commit-point manually
 - check that DB engine supports delayed commit

```
con.setAutoCommit(false);    // turn off auto-commit
con.commit();                // commit manually
con.setAutoCommit(true);     // turn on again
```

Manual Commit

```
con.setAutoCommit(false);  
...  
Statement stmt = null;  
try {  
    stmt = con.createStatement();  
    int count = stmt.executeUpdate(  
        "UPDATE lecturer SET office = 119"  
        + " WHERE name = 'A Simons'");  
    con.commit();  
}  
catch (SQLException ex) {  
    ex.printStackTrace();  
}  
finally {  
    if (stmt != null) stmt.close();  
}
```

Updates are only committed at the manual commit-point.

Similar behaviour for all updates using this open Connection, unless you switch back



Transactions

- Transaction
 - a database **transaction** is a single, complete unit of work, which must either execute **completely**, or **not at all**
 - must be **ACID**: **A**tomic (all or nothing), **C**onsistent (data integrity), **I**solated (serializable), **D**urable (permanent)
- Protect against brief loss of service
 - wrap a set of updates in a **transaction**, if all updates must happen together (e.g. credit/debit in a money transfer)
 - the transaction must succeed, or fail, as a whole
 - upon failure, the database must **rollback** (forget temporary changes) or **revert** (recover the “**before image**” of the data)

Transaction Support

- MySQL support for transactions is engine-dependent
 - **InnoDB** engine (the default) supports ACID transactions
 - **MyIASM** engine only has atomic actions (auto commit=true)
 - **BDB** (Berkley DB) engine also supports transactions
- Can specify storage engine when defining a table
 - check this engine is available for your DB
 - use the ENGINE keyword in MySQL:

```
CREATE TABLE account(  
    accno INT NOT NULL PRIMARY KEY AUTO_INCREMENT,  
    holder VARCHAR(50),  
    balance INT  
) ENGINE=BDB;
```

or:

```
CREATE TABLE lecturer  
(...) ENGINE=InnoDB;
```

Transactions in JDBC

```
con.setAutoCommit(false);           // turn off auto-commit
Statement credit = null;             // SQL statement objects
Statement debit = null;
try {
    credit = con.createStatement();
    debit = con.createStatement();
    debit.executeUpdate("UPDATE account ...");
    credit.executeUpdate("UPDATE account ...");
    con.commit();                     // manually commit when ready
}
catch (SQLException ex) {
    if (con != null) con.rollback(); // if transaction fails
}
finally {
    if (credit != null) credit.close();
    if (debit != null) debit.close();
}
```



SQL Injection!

- Control access to data
 - ensure users have restricted views of data
 - ensure users have appropriate authorisation
- Validate all inputs
 - biggest mistake is **failure to validate** all inputs
 - allows **SQL injection faults**, could kill the DB
- Encrypt the data
 - if you really must!

injecting an extra SQL command
in a simple text entry field

Enter student ID:

0011234567; drop table
student

Prepared Statements

- For common queries with similar patterns
- Pre-compile a prepared statement with ? parameters
- Protects against injection: actual values \neq SQL commands

```
// prepare beforehand
PreparedStatement pstmt = con.prepareStatement(
    "SELECT id, name FROM lecturer WHERE name=?");

// use later

pstmt.setString(1, "A Simons");
ResultSet res = pstmt.executeQuery();
...
pstmt.close();
```

A parameter in the pstmt
is shown as ?

setX(nth, val) sets the
nth param with val

Repeated Actions

```
// prepared statement with two parameters
// column 1 = null because of auto-increment

PreparedStatement pstmt = con.prepareStatement(
    "INSERT INTO lecturer VALUES (null, ?, ?)");

// iterate over a list of lecturers in memory and
// insert these into the database

for (Lecturer lect : department.getLecturers()) {
    pstmt.clearParameters();
    pstmt.setString(1, lect.getName());
    pstmt.setInt(2, lect.getOffice());
    int count = pstmt.executeUpdate();
}
...
pstmt.close();
```

Each parameter in the pstmt is shown as ?

clearParameters() clears all old values

setX(nth, val) sets the nth param with val



Concurrency Control

- Default strategy
 - JDBC uses the default strategy for the underlying DB
- Custom strategy
 - get and set using methods of the `Connection` class

```
int level = con.getTransactionIsolation();           // find the default setting
con.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE);
                                                    // serialize transactions
```

Constants for levels of concurrency control:

```
TRANSACTION_NONE, TRANSACTION_READ_UNCOMMITTED,
TRANSACTION_READ_COMMITTED, TRANSACTION_REPEATABLE_READ,
TRANSACTION_SERIALIZABLE
```



Lab 2: Input Validation

Run a Poll

- Sketch the Java for your input validation
 - assume you have a Java Swing `TextField queryField`;
 - this has a `getText()` method that returns the text that was entered into the `queryField`, as a `String`
 - you need to determine whether this `String` is healthy
- How will you check the input?
 - what `String` API functions can you use to look for SQL-injection attempts?
 - should you raise an exception if an attack is being made?
 - can you clean up the text `String` from the `queryField`?



A decorative graphic consisting of overlapping yellow, red, and blue squares with a black crosshair.

MySQL Accounts

- Issuing group accounts
 - MySQL accounts are generated annually for group-work by the DCS support team
 - they send a list of new accounts to your lecturer who then has to allocate to team-leaders
 - team leaders watch out for this in your email inbox!
- Interpreting login data
 - the email will mention a DBname and a DBpassword
 - e.g. DBname=team043 and DBpassword=5a94128d
 - note DBname is used also as your team's userID.
 - need a DBname, userID and DBpassword to connect

A decorative graphic consisting of overlapping yellow, red, and blue squares with a black crosshair.

MySQL Client

- MySQL command-line client (shell, monitor)
 - be connected to the University's VPN
 - in Linux, open a terminal or console

```
% mysql -hstusql.dcs.shef.ac.uk -uteam043 -p team043
```

```
Enter password: *****)
```

```
Welcome to the MySQL monitor. ...
```

- -h specifies the host, the database server
- -u specifies the userid, your team name
- -p prompts for your password on the next line (why?)
- the last argument is the database name (why?)

Using the Client

- Type any SQL instructions at the prompt

```
% mysql -hstusql.dcs.shef.ac.uk -uteam043 -p team043
```

```
Enter password: *****
```

```
Welcome to the MySQL monitor.  Commands end with ; or \g
```

```
...
```

```
mysql> CREATE TABLE Student (  
-> regno INTEGER NOT NULL PRIMARY KEY,  
-> forename VARCHAR(30),  
-> surname VARCHAR(30));
```

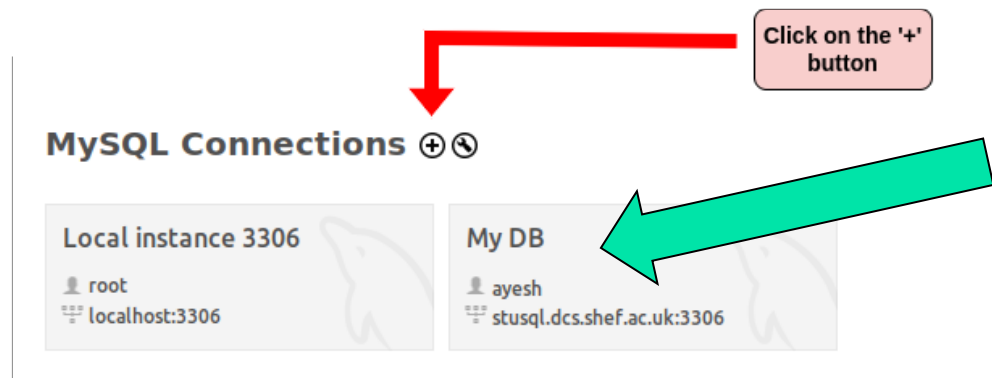
```
Query OK, 0 rows affected (0.42 sec)
```

```
mysql>
```

prompts with -> until the
command is terminated

MySQL Workbench

- MySQL GUI-based client, for Windows (see Lab for Linux)
 - <https://dev.mysql.com/downloads/workbench/>
 - Version: MySQL Workbench 8.0.34
 - Remember, connection to DCS servers requires VPN first

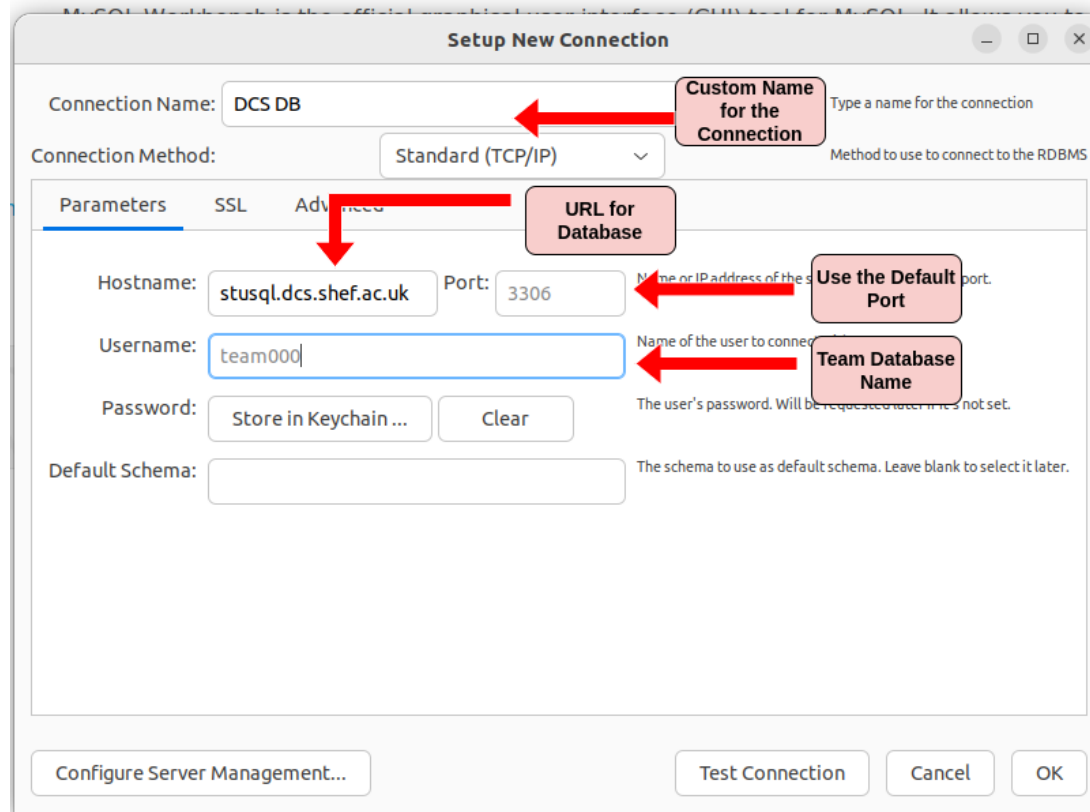


The screenshot shows the 'MySQL Connections' pane in MySQL Workbench. It contains two connection entries: 'Local instance 3306' (with user 'root' and host 'localhost:3306') and 'My DB' (with user 'ayesh' and host 'stusql.dcs.shef.ac.uk:3306'). A red arrow points to the '+' icon next to the 'MySQL Connections' header, with a callout box saying 'Click on the '+' button'. A green arrow points to the 'My DB' connection entry.

The first time you use, you won't see the list of MySQL connections shown here.

But click on (+) to get the following...

Configuring



Setup New Connection

Connection Name: DCS DB

Connection Method: Standard (TCP/IP)

Parameters | SSL | Advanced

Hostname: stusql.dcs.shef.ac.uk Port: 3306

Username: team000

Password: Store in Keychain... Clear

Default Schema:

Buttons: Configure Server Management..., Test Connection, Cancel, OK

Give the connection any name you like

Use the standard TCP/IP

Host is stusql.dcs.shef.ac.uk
Port is 3306

Username is your team DB name

Security



Username: Name c

Password: Store in Keychain ... Clear The use

Click on "Store in Keychain"

Store Password For Connection

Please enter password for the following service:

Service: Mysql@stusql.ics.shef.ac.uk:3306

User: team000

Password:

OK Cancel

Enter the Password

Setup New Connection

Connection Name: Type a name for the connection

Connection Method: Method to use to connect to the RDBMS

Parameters SSL Advanced

Use SSL If available

SSL Key File: Require Path to Client Key File for SSL.

SSL CERT File: Require and Verify CA Path to Client Certificate File for SSL.

SSL CA File: Require and Verify Identity Path to Certificate Authority file for SSL.

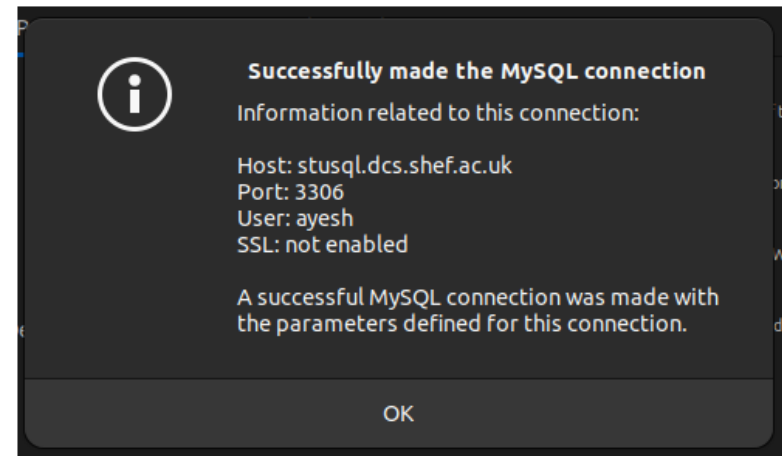
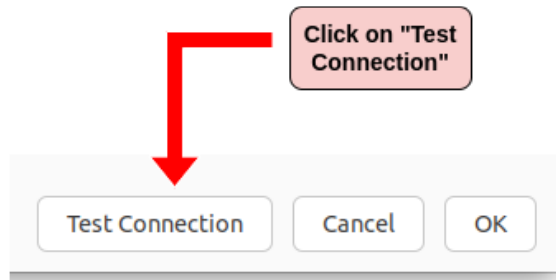
Click on Use SSL -> No

Click on "store in keychain" to enter team password


Click on SSL tab and select No

(stustore does not support SSL connections)

Connecting



MySQL Connections ⊕ ⊞



Click on Test Connection to finish set-up

Next time, just click on the button for the configured DB



Summary

- Java and databases have different strengths of type
- MySQL is the best free database to use on Windows and Linux
- JDBC is the bridge between Java and any database, using a suitable driver which is easy to install
- Connections, statements, result sets must **always be closed** using **auto-closing** resource management or the **try-catch-finally** idiom
- Prepared statements, manual commit and robust transactions are possible with some DB engines
- Concurrency control is possible, but default settings are OK
- Guard against **SQL injection** through input validation!