

Systems Design and Security



Part 10: Security and Robustness

<http://staffwww.dcs.shef.ac.uk/people/A.Simons/>

Home \Rightarrow Teaching \Rightarrow Lectures
 \Rightarrow COM2008/COM3008



Bibliography



- Software Engineering
 - I Sommerville. Software Engineering, 10th ed., Pearson, 2022.
- Database Systems
 - T Connolly and C Begg, Database Systems – a Practical Approach to Design, Implementation and Management, 6th ed., Pearson, 2014.
- Security Engineering
 - R Anderson. Security Engineering, 3rd ed., John Wiley, 2022. <https://www.cl.cam.ac.uk/~rja14/book.html>



Outline


- Authentication and authorisation
- Integrity and non-repudiation
- Confidentiality and encryption
- Robustness via redundancy
- Restricted access to views
- Penetration resistance
- Availability via concurrency

Reading: Sommerville chapters 10-14;
Connolly and Begg, chapters 20, 22, 24-26;
Anderson – the whole book!

What is Security?

Combination of
safety and trust

Four main pillars



<http://www.freevector.com/jigsaw-puzzle>

Authentication

The identities of all parties concerned should be verified

This includes both users and software

Authentication

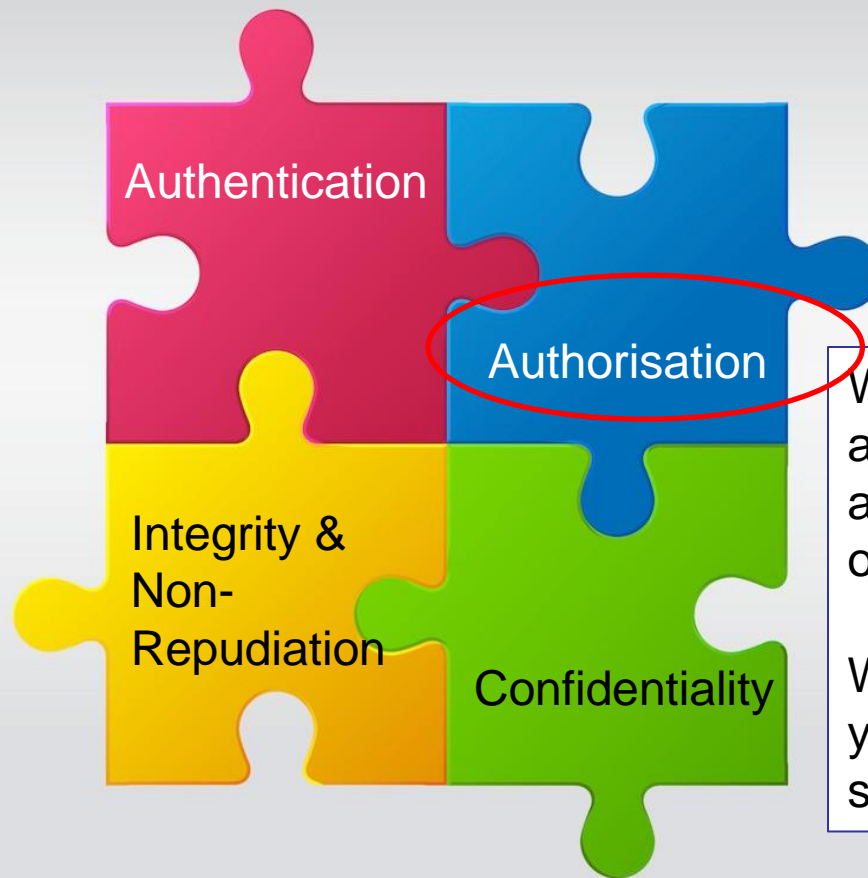
Authorisation

Integrity &
Non-Repudiation

Confidentiality

<http://www.freevector.com/jigsaw-puzzle>

Authorisation



What operations are user-roles allowed to carry out?

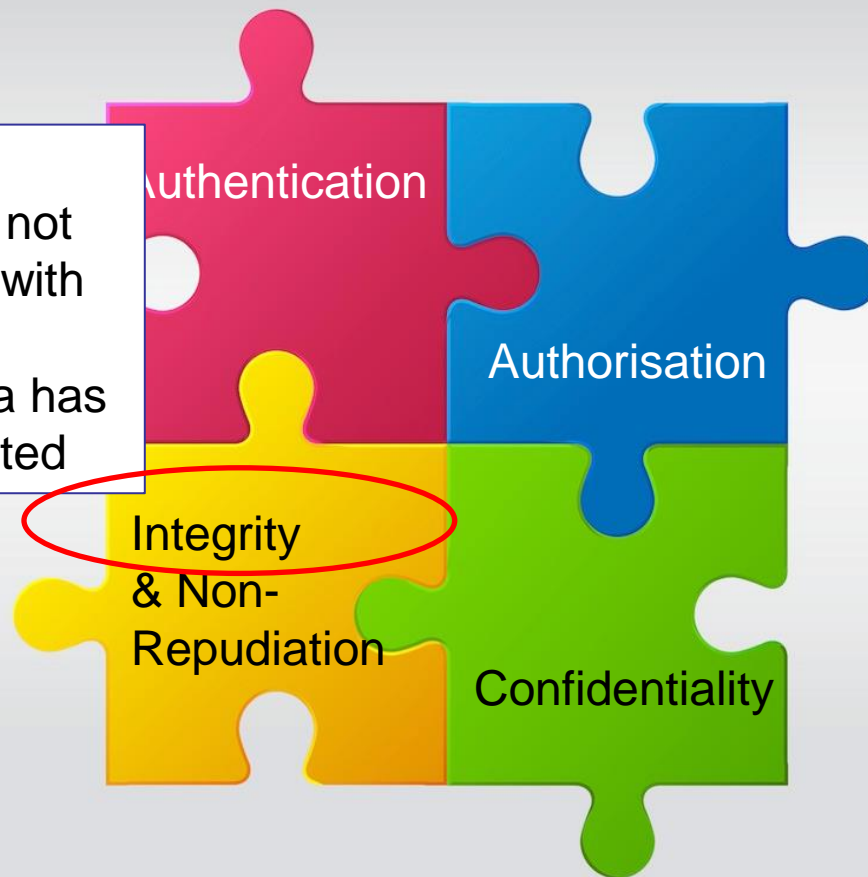
What access to your system is the software granted?

<http://www.freevector.com/jigsaw-puzzle>

Integrity

Ensure that the information has not been tampered with

Ensure that data has not been corrupted

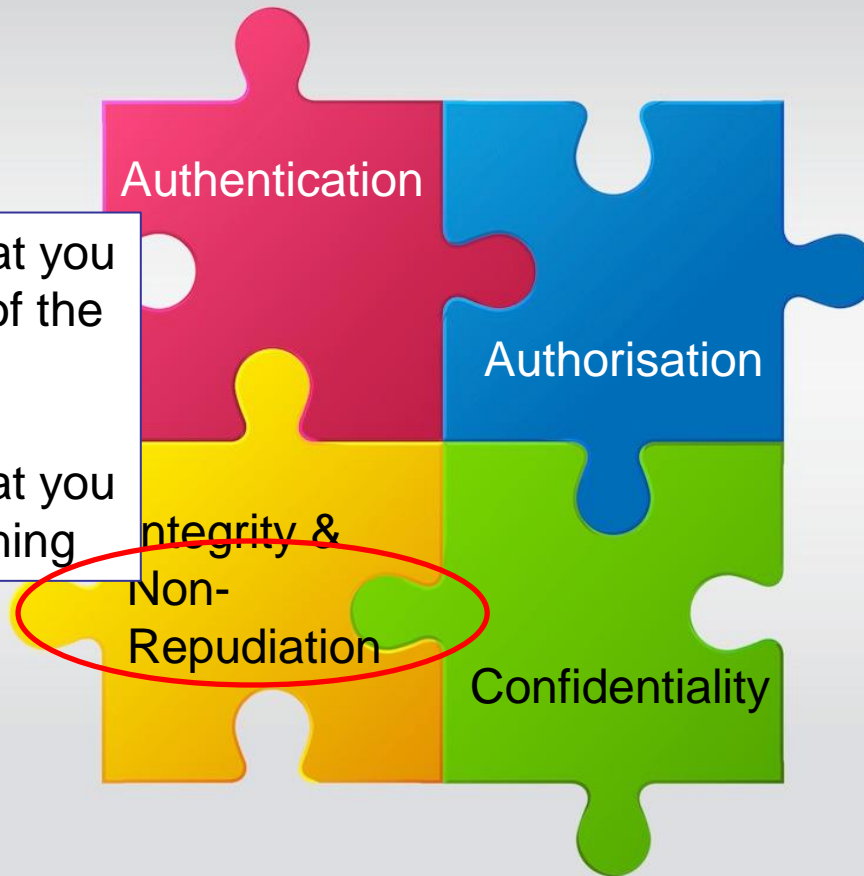


<http://www.freevector.com/jigsaw-puzzle>

Non-Repudiation

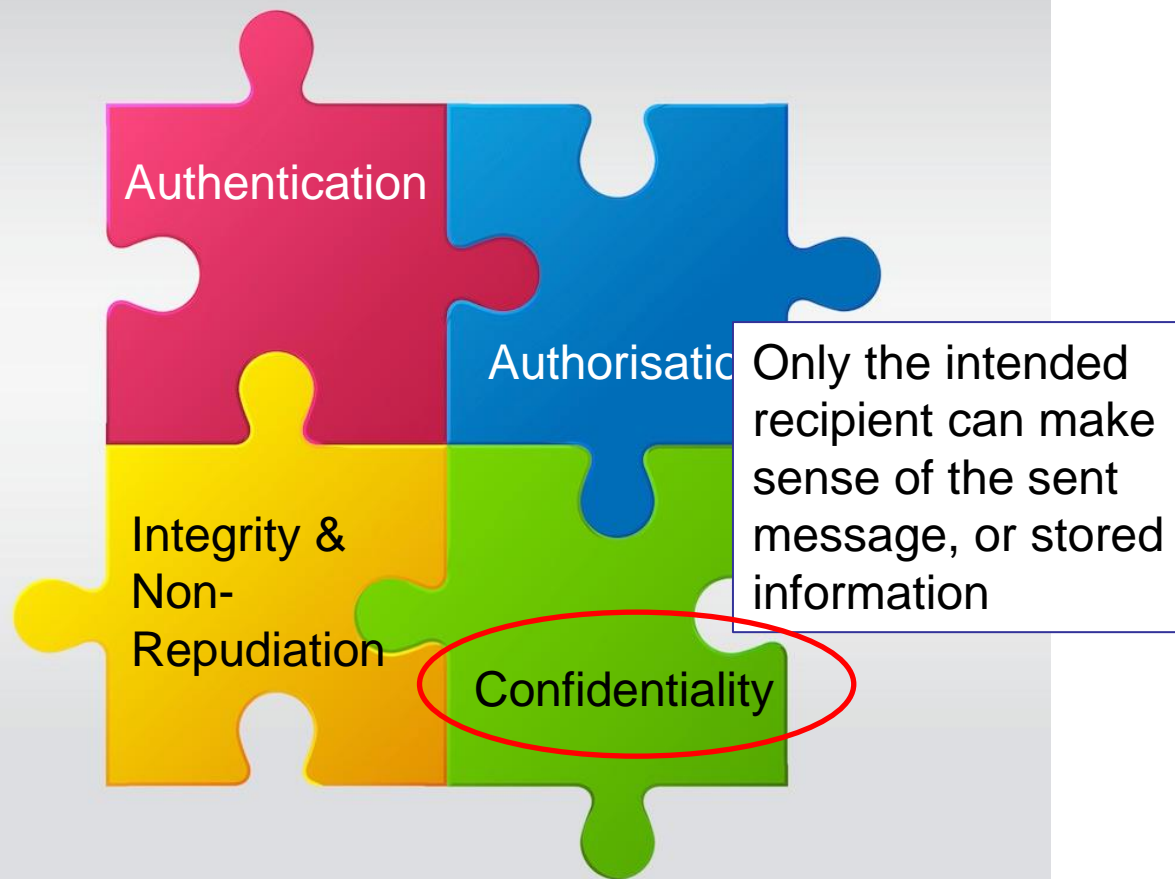
Cannot deny that you
are the sender of the
information

Cannot deny that you
received something



<http://www.freevector.com/jigsaw-puzzle>

Confidentiality



<http://www.freevector.com/jigsaw-puzzle>



Tradeoffs

- With unlimited resources, most forms of security can eventually be broken
 - arms race between the criminal and the security expert
- Cost of breaking must outweigh the reward
 - make it harder for them to crack your system (compared to others)
- Must consider end-to-end security
 - only as secure as the weakest link
 - eg: you have private key encryption, but how secure is your key store?
- Aim for simplicity
 - you don't want to frighten away legitimate users
 - security regime must be easy to maintain



Security Scenarios

- Online banking
 - **authentication**: is this a valid account holder?
 - **authorisation**: do they have permission to access this account?
 - **confidentiality**: is the account data secure?
 - is security vetting easy for legitimate users?
- Downloading software
 - **authentication**: does the code come from a trusted source?
 - **integrity**: has the code been tampered with, before or during downloading?
 - **authorisation**: does the code have permission to install itself on your machine, or remove other files?
 - will the software only do what it claims to do?



Security Scenarios

- Online credit transactions
 - **authentication:**
 - does the credit card belong to the customer?
 - is the merchant valid?
 - is the merchant bank valid?
 - **integrity:**
 - have any details been altered *en route*?
 - **non-repudiation:**
 - can the purchaser deny that they ordered the item?
 - can the vendor deny that they received the money?
 - **confidentiality:**
 - should the merchant have access to your credit card?
 - should the bank have access to the purchase details?



Authentication

- Password-controlled access
 - every user given a unique loginID
 - every user given a secret password
 - password of 8+ chars, mixed case, digits and symbols
- Pre-registration schemes
 - e.g. HTTP htaccess: 64-bit encrypted user:password
 - e.g. HTTP digest: better encryption of user-info
 - server decodes against a fixed list of known users
- Self-registration schemes
 - form-based self-registration, dynamic list of known users
 - requires secure transmission, storage of passwords
 - e.g. encrypted with additional "salt" data in DB
 - requires super-user determination of role-based access (next)



Authorisation

- Mandatory Access Control (MAC)
 - controlled by a security manager in high-security systems
 - e.g. permission for a thread to touch certain data
- Discretionary Access Control (DAC)
 - permissions set as desired by the data owner
 - e.g. Unix user, group and world access rights to files
- Role-based Access Control (RBAC)
 - certain users are granted certain roles
 - each role is granted certain permissions
 - these concern different target objects (files, ports)
- Lattice-based Access Control (LBAC)
 - a user must exceed access level of protected accessible object
 - multiple users granted **greatest lower bound** (meet) access
 - multiple objects offer **least upper bound** (join) access



Role-Based

- Who or what?
 - individual users, or groups
 - role-based access, eg: manager, customer
 - specific programs, from a given origin
 - Which resource?
 - individual files, DB tables
 - files of certain template-types
 - files from certain locations
 - What restriction?
 - read only
 - read and write
 - update or remove files
 - connect, listen to channel
- Bespoke methods grant **role-based permissions to users**
 - Java also has features to support **authorisation for software**



Restricted Users

- SQL privileges grant selective access to different users

```
GRANT SELECT ON MyDB.Borrower  
TO PUBLIC;
```

allows anyone to select on the
Borrower table in MyDB

```
CREATE USER 'user1'@'localhost'  
IDENTIFIED BY 'passwd1';
```

creates user1 with passwd1

```
GRANT ALL ON MyDB.Loan  
TO 'user1'@'localhost';
```

allows user1 to insert, delete,
select, update the Loan table

```
GRANT SELECT, UPDATE(issueDate, dueDate)  
ON MyDB.Loan TO 'user1'@'localhost';
```

allows user1 to select, update two
columns of the Loan table

```
GRANT UPDATE ON MyDB.BookTitle TO  
'user1'@'localhost' WITH GRANT OPTION;
```

allows user1 to update BookTitle;
and to pass on this privilege to
others



Restricted Roles

- SQL roles grant groups of privileges to groups of users

```
CREATE ROLE 'SuperUser', 'User';
```

defines two roles named SuperUser and User

```
GRANT ALL ON MyDB.* TO 'SuperUser';
```

allows SuperUser role to have all permissions on MyDB

```
GRANT SELECT ON MyDB.* to 'User';
```

allows User role only to have viewing permissions

```
GRANT 'SuperUser' TO 'user1'@'localhost';
```

```
GRANT 'User' TO 'user2'@'localhost';
```

```
GRANT 'User' TO 'user3'@'localhost';
```

grants roles with the given permissions to individual users

```
REVOKE 'User' FROM 'user2'@'localhost';
```

revoke roles and permissions



Lab 1: User Accounts

Run a Poll

- How will your database set up user accounts?
 - your MySQL DB has only one secure login for the team, with a team password (secret in your Java code)
 - you cannot create extra MySQL accounts on stustore
 - does your application support self-registration?
 - how will you handle setting up new user accounts?
 - how will new users be assigned suitable privileges?
- What tables or Java code will you need:
 - to handle usernames and passwords?
 - to offer role-based access permissions?



Integrity

- Secure transmission
 - authentication: comes from a valid source
 - integrity: has not been modified/corrupted in transit
- Main technologies
 - authentication: digital signatures, certificates
 - integrity: message digests (digital fingerprints)
- Both use encryption
 - Secure Hash Algorithm (SHA) for one-way encoding of message digests - evidence of no tampering
 - shared key encryption (but problem of key sharing, storage)
 - Rivest-Shamir-Adleman (RSA) algorithm for public/private key encryption - avoids problem of key storage and sharing

Cryptographic Hash

"I am a secret
message"



A45D761F0EBB271C9BC

"I am a secret
message"



F39E76CCE039882645D8

- keyless algorithm, easy to compute
- fixed-length digests, 100-200 bits
- hard to recover the clear message
- hard to obtain a collision, $h(m1) = h(m2)$
- discontinuous mapping, $h(m) \neq h(m')$



Secure Hash Algorithm

- SHA-1 from NSA and NIST
 - block size 512 bits
- generates 160-bit hash
 - 2^{160} digests
- previous competitors
 - MD4 (weak)
 - MD5 128-bit hash
- how secure?
 - attack found in 2005
 - from 2010, deprecated
 - SHA-2, SHA-3 now recommended

```
int a, b, c, d, e = ...

for (t=0; t < 79; ++t) {
    int tmp = c << 5 +
        f(b,c,d) + e + W[t]
        + K[t];
    e = d;
    c = b << 30;
    b = a;
    a = tmp;
}

// arithmetic is mod 32
```



Message Authentication

- Simple message digest
- sender sends:
 - message m
 - digest $d = h(m)$
- receiver computes
 - digest $d' = h(m)$
 - checks $d = d'$
- weakness
 - if both d , m modified in transit
- MAC: Message authentication code
 - shared password p
- sender computes:
 - digest $d = h(p || m)$
 - $mac = h(p || d)$
 - sends m , mac
- receiver computes
 - digest $d' = h(p || m)$
 - checks $mac = h(p || d')$



Confidentiality

- Symmetric key encryption
 - use same key to encrypt and decrypt
 - or determine one key from the other
 - problem with secure key sharing, storage
 - needs new key for each pair of parties
 - key management problem
- Public/private key encryption
 - uses two related prime keys, very large
 - impossible to guess one from the other
 - sender encrypts with receiver's public key
 - receiver decrypts with own private key
 - no key transmission issues

Encryption / Decryption

"original message"



shared key



10110101110101



shared key



decrypt

"recovered message"



"original message"



encrypt



10110101110101



decrypt


private key



"recovered message"




public key



RSA Algorithm

Ron Rivest, Adi Shamir, Len Adleman





- Choose n , the product of two large primes p and q
- Choose e and d with very specific mathematical properties in relation to n
 - e.g. $\gcd(d, p-1)=1$, $\gcd(e, q-1)=1$
- e and n form the public key, released
- d is the private key, never released
- message m , is encrypted to cipher c , as
 - $c = m^e \bmod n$
- cypher c is decrypted, yielding message m , as
 - $m = c^d \bmod n$

<http://www.usc.edu/dept/molecular-science/RSA-2003.htm>



Cryptanalysis

- 1974 Data Encryption Standard (IBM)
 - 56-bit key cracked in 22hrs in 1999 using 100,000 computers; decertified; is probably easily cracked now
 - 64-bit keys breakable by governments, criminal organisations
 - 128-bit keys good for foreseeable future (unless quantum comp)
-  ■ 1976 Diffie & Hellmann public key cryptography
 - prime factors of large products ($> 10^{100}$) very hard to find, but need large keys > 512 bits
 - Sarah Flannery's algorithm speeds up processing x30 in 1999
-  ■ 1977 RSA (Rivest-Shamir-Adleman) public/private key algorithm
 - actually GCHQ spy Clifford Cocks had this from 1973, secret till 1997!
 - 512-bit keys can be cracked in weeks using dedicated machines
 - 1024-bit, 2048-bit keys good for many decades?



Signatures, Certificates

- Digital Signature
 - combination of message digest and encryption
 - sender encrypts digest with own private key
 - receiver decrypts digest with sender's public key
 - authenticates the sender, and message integrity
- Digital Certificate
 - certifying authority (CA) vouches for third party by issuing certificate with public key
 - first party trusts CA, so uses this public key to authenticate third party
 - example CAs: Verisign, Entrust



Java Security

- Java Sandbox
 - a "safe environment" in which to run Java code
 - can prevent reading/writing to files on host machine
 - originally for applets; now for any Java application
 - uses a security manager (disabled by default)
- Security elements
 - **permissions**: actions that code is allowed to perform
 - **code sources**: origin of code, digital signatures
 - **protection domains**: dictionary that maps permissions to code sources
 - **key stores**: digital signatures use keys, held in store
 - **policy files**: list permissions, identify keystore location, specify code sources and protection domains



Java Policy File

```
keystore "${user.home}${/}.keystore";
```

```
grant codeBase "http://www.dcs.sheffield.ac.uk/" {  
    permission java.io.FilePermission "/tmp", "read";  
    permission java.lang.RuntimePermission "queuePrintJob";  
};  
grant signedBy "ajhs" codeBase "http://ajhs.com/" {  
    permission java.security.AllPermission;  
};  
grant signedBy "gjb" {  
    permission java.net.SocketPermission " *:1024-",  
        "accept, connect, listen";  
};  
grant {  
    permission java.util.PropertyPermission  
        "java.version", "read";  
};
```

code from a
given origin

a type of
permission

Protection Domains

```
keystore "${user.home}${/}.keystore
```

allow code loaded from DCS to read /tmp and to queue print jobs

```
grant codeBase "http://www.dcs.sheffield.ac.uk/" {  
    permission java.io.FilePermission "/tmp", "read";  
    permission java.lang.RuntimePermission "queuePrintJob";  
};
```

```
grant signedBy "ajhs" codeBase "http://ajhs.com/" {  
    permission java.security.AllPermission;  
};
```

allow code loaded from ajhs.com and signed by ajhs to do anything

```
grant signedBy "gjb" {  
    permission java.net.SocketPermission  
        "accept, connect, listen";  
};
```

allow code signed by gjb to do 3 actions on all hosts, ports > 1024

```
grant {  
    permission java.util.PropertyPermission  
        "java.version", "read";  
};
```

allow all code to read the Java version



Java Example

```
import java.net.*;
import java.security.*;

class AccessTest {
    public static void main(String[] args) {
        SocketPermission sp = new SocketPermission(
            "my.host.name:6000", "connect");
        try {
            AccessController.checkPermission(sp);
            System.out.println("OK to open socket");
            // carry on here ...
        }
        catch (AccessControlException ace) {
            ace.printStackTrace();
        }
    }
}
```


create a
permission

test the
permission

exit if
refused

extend `java.security.BasicPermission` to
create your own kinds of permission

What is Robustness?

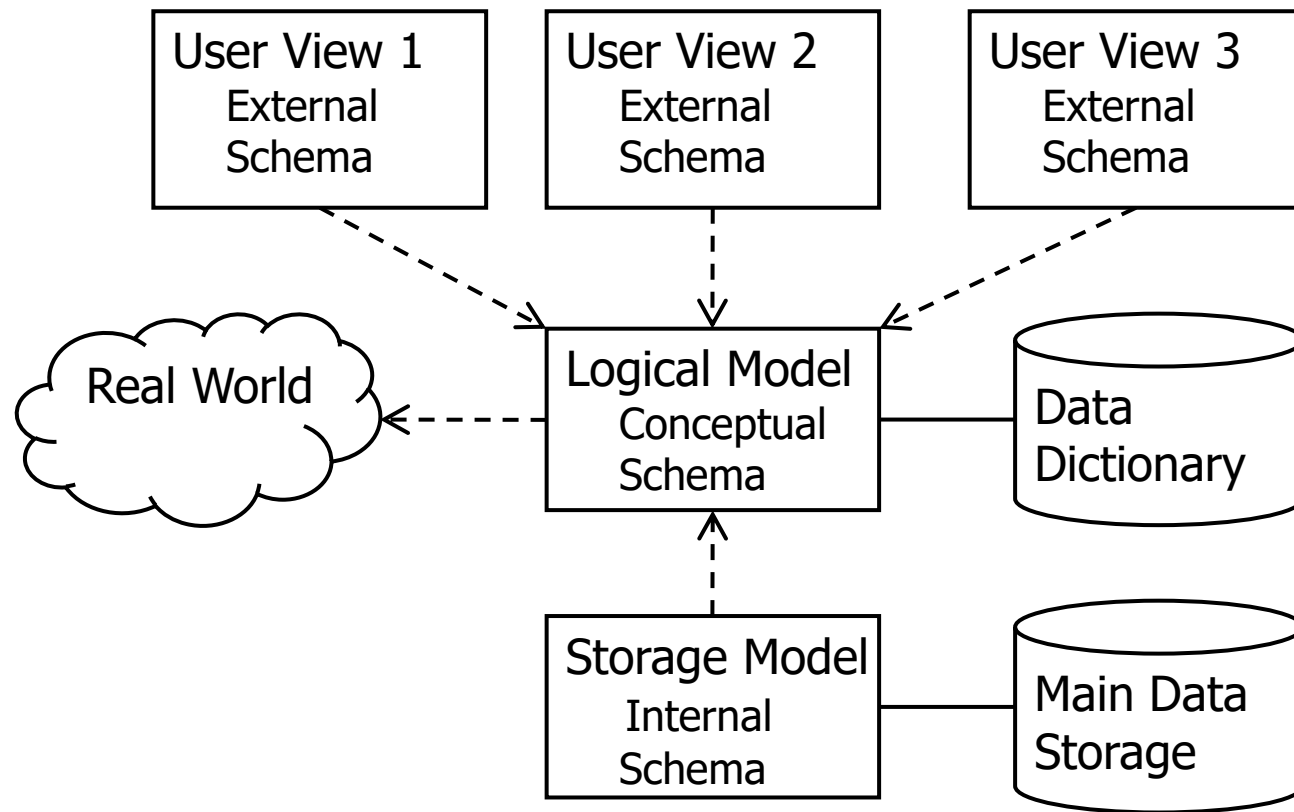




Physical Protection

- Protect against loss of main data centre
 - multiple copies of database in **separate locations**
 - but expensive to update live data simultaneously
 - **regular backups** of whole dataset to remote locations
 - data storage in “**the cloud**” – no more big office servers?
 - trust issues – who else can see data in the cloud?
- Protect against local hardware failure
 - Use **RAID** architecture (**R**edundant **A**rray of **I**ndependent **D**isks)
 - RAID-0: data is “striped” (segmented) across different disks
 - RAID-1: mirror copy of data is also made, across different disks
 - RAID-2: error-correcting codes also retained, on different disks

Controlled Views





View Definition

- A view is a derived table, computed from other tables
- Views enhance usability, security in large, complex databases
- Can be inefficient to query – hides complex subquery


```
CREATE VIEW Overdue (memberID, forename, surname, overdue) AS
SELECT memberID, forename, surname, COUNT (*) AS overdue
FROM Borrower, Loan WHERE dueDate < '2022-10-01'
GROUP BY memberID, forename, surname
```

----- must name all non-aggregate columns in the grouping clause

Overdue

memberID	forename	surname	overdue
1012234	John	Smith	4
1667753	Jane	Doe	1
1784532	Jaswinder	Singh	3

UML Profile for View



views are only updatable if they have a **single** base table and include the same PK

- UML Profiles
 - UML database profile is not universal; but widespread
- View tables
 - use the «view» label to indicate a view **table**, and
 - {read only} constraint if the view **cannot** be updated
- Any key attributes?
 - only if using same PK as the **single** base table
 - candidate keys are not reliable after 'project'



Penetration Resistant

- Control access to data
 - ensure users have restricted views of data
 - prohibit arbitrary free-data entry where possible
 - provide restricted-choice selections instead
- Don't forget: validate all inputs
 - #1 security error is **failure to validate** all inputs
 - **SQL injection faults** could kill the DB

injecting an extra SQL command
in a simple text entry field

Enter student ID:

0011234567; drop table
student



Lab 2: Read-only View



- Write the SQL for a view of the marks obtained by a student for studied modules
 - assume table Student (regNo, forename, surname)
 - assume table Module (modID, modName, lecturer)
 - assume table Study (regNo, modID, mark1, mark2)
- Outline syntax for a view
 - CREATE VIEW ... AS SELECT ... FROM ... WHERE ...;
 - no GROUP BY as we are not using any aggregate functions like COUNT(*) here
- We want to see in the resulting view
 - regNo, forename, surname, modID, modName, mark1, mark2



Updatable View?

Run a Poll

- Is this view updatable?

```
CREATE VIEW Transcript (regNo, surname, forename,  
modID, modName, mark1, mark2) AS  
  SELECT regNo, surname, forename, modID, modName,  
    mark1, mark2 FROM (Student, Module, Study)  
  WHERE Student.regNo = Study.regNo  
    AND Module.modID = Study.modID;
```



Failsafe Transactions

- Transaction
 - a database **transaction** is a single, complete unit of work, which must either execute **completely**, or **not at all**
 - must be **ACID**: **A**tomic (all or nothing), **C**onsistent (data integrity), **I**solated (serializable), **D**urable (permanent)
- Protect against brief loss of service
 - wrap a set of updates in a **transaction**, if all updates must happen together (e.g. credit/debit in a money transfer)
 - the transaction must succeed, or fail, as a whole
 - upon failure, the database must **rollback** (forget temporary changes) or **revert** (recover the “**before image**” of the data)



Transaction Recovery

- Commit and rollback, or revert?
 - **commit**: make temporary changes permanent
 - **rollback**: discard temporary changes, before a commit
 - **revert**: restore DB to earlier state from saved old data
- Three possible strategies
 - **deferred update**: write changes to a temporary log file, commit all changes together, or **rollback** (discard) if incomplete (DB may have uncommitted changes)
 - **immediate update**: log the old value, change the main data, recover by **reverting** (extra update, but DB is consistent)
 - **shadow paging**: keep dynamic pointers to blocks of current, old data and **merge** or **revert** (needs memory management)



Concurrent Availability

- Concurrent access
 - multiple users access the same database simultaneously
 - transactions overlap, complete in nondeterministic order
- Concurrency conflicts
 - **dirty read** – user1 reads data, as user2 is updating it, but has not committed; the data read is **out of date**
 - **non-repeatable read** – user1 runs a query twice, as user2 is modifying data; the query returns **inconsistent data**
 - **phantom read** – user1 runs a query twice, as user2 is adding to the data; the query returns extra **phantom data**
 - **lost update** – user1 and user2 make simultaneous updates to the same data; the last to commit is saved, the first **update is lost**



Concurrency Control - I

- Serialize all transactions
 - pro: all transactions forced to execute in **sequential order**
 - con: serial bottleneck reduces performance
- Row/page locking
 - pro: other transactions **locked out** while row is being updated
 - con: possible deadlock if transactions wait for each other
- Deadlock avoidance
 - pro: aborts transactions if all locks **cannot be obtained**
 - con: aborts and restarts too many valid transactions
- Deadlock detection
 - uses a **wait-for graph** to detect and **break circular deadlocks**, aborts an arbitrary transaction; results in fewer aborts



Concurrency Control - II

- Shared and exclusive locks
 - **shared** read-lock allows multi-user access for reading
 - **exclusive** write-lock prevents all further access while updating
- Phantom and intent locks
 - **phantom lock** reserves an empty row for later insertion
 - **intent lock** declares intention to commit changes, may co-exist with shared locks; becomes **exclusive** only during the commit phase
- Timestamping
 - each transaction is given a unique start-time and all rows are marked with their **last-read** and **last-written** timestamps
 - transactions can freely **write** any rows with **earlier** read/write stamps, and **read** rows with **later** read stamps; otherwise abort



Summary

- Authentication determines that a user is who they claim to be
- Authorisation determines what functions or data a user, or other software, has the right to use, access or update
- Integrity determines whether data has been tampered with during transit and whether it comes from a trusted source
- Confidentiality ensures that only the agreed parties can see data secured by encryption (shared key, or public-private key)
- Physical attacks and equipment failure are mitigated by distribution and data-replication
- Penetration attacks are mitigated by offering restricted views of data and strong input validation
- Operational failures are mitigated by transaction recovery
- Availability is assured partly through concurrent access